

podkohlxo

April 21, 2024

1 Worksheet 21

Name: JINGYI ZHANG UID: U26578499

1.0.1 Topics

- Logistic Regression
- Gradient Descent

1.1 Logistic Regression

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import sklearn.datasets as datasets
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import PolynomialFeatures

centers = [[0, 0]]
t, _ = datasets.make_blobs(n_samples=750, centers=centers, cluster_std=1,
    ↪random_state=0)

# LINE
def generate_line_data():
    # create some space between the classes
    X = np.array(list(filter(lambda x : x[0] - x[1] < -.5 or x[0] - x[1] > .5,
    ↪t)))
    Y = np.array([1 if x[0] - x[1] >= 0 else 0 for x in X])
    return X, Y

# CIRCLE
def generate_circle_data(t):
    # create some space between the classes
    X = np.array(list(filter(lambda x : (x[0] - centers[0][0])**2 + (x[1] -
    ↪centers[0][1])**2 < 1 or (x[0] - centers[0][0])**2 + (x[1] -
    ↪centers[0][1])**2 > 1.5, t)))
    Y = np.array([1 if (x[0] - centers[0][0])**2 + (x[1] - centers[0][1])**2 >=
    ↪1 else 0 for x in X])
```

```

    return X, Y

# XOR
def generate_xor_data():
    X = np.array([
        [0,0],
        [0,1],
        [1,0],
        [1,1]])
    Y = np.array([x[0]^x[1] for x in X])
    return X, Y

```


a) Using the above code, generate and plot data that is linearly separable.

```

[2]: X, Y = generate_line_data()

plt.figure(figsize=(8, 6))
plt.scatter(X[Y == 0][:, 0], X[Y == 0][:, 1], color='blue', label='Class 0')
plt.scatter(X[Y == 1][:, 0], X[Y == 1][:, 1], color='red', label='Class 1')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.grid(True)
plt.show()

```



b) Fit a logistic regression model to the data and print out the coefficients.

```
[3]: model = LogisticRegression().fit(X, Y)
      model.coef_, model.intercept_
```

```
[3]: (array([[ 4.11128306, -4.10408124]]), array([0.06146435]))
```

c) Using the coefficients, plot the line through the scatter plot you created in a). (Note: you need to do some math to get the line in the right form)


```
[4]: plt.figure(figsize=(8, 6))
      plt.scatter(X[Y == 0][:, 0], X[Y == 0][:, 1], color='blue', label='Class 0')
      plt.scatter(X[Y == 1][:, 0], X[Y == 1][:, 1], color='red', label='Class 1')

      feature1 = np.array([X[:, 0].min(), X[:, 0].max()])
      feature2 = -(model.intercept_[0] + model.coef_[0][0] * feature1) / model.
                  ↪coef_[0][1]

      plt.plot(feature1, feature2, label='Decision Boundary')

      plt.xlabel('Feature 1')
```


```
plt.ylabel('Feature 2')
plt.legend()
plt.grid(True)
plt.show()
```



d) Using the above code, generate and plot the CIRCLE data.

```
[5]: X, Y = generate_circle_data(t)

plt.figure(figsize=(8, 6))
plt.scatter(X[Y == 0][:, 0], X[Y == 0][:, 1], color='blue', label='Class 0')
plt.scatter(X[Y == 1][:, 0], X[Y == 1][:, 1], color='red', label='Class 1')
plt.title('Circle Data')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.grid(True)
plt.show()
```



e) Notice that the equation of an ellipse is of the form

$$ax^2 + by^2 = c$$

Fit a logistic regression model to an appropriate transformation of X.

```
[6]: #model = ...
from sklearn.preprocessing import PolynomialFeatures

# Redefine polynomial features to only include squares of each feature
poly_square = PolynomialFeatures(degree=2, include_bias=False,
    ↪interaction_only=False)
X_circle_square = poly_square.fit_transform(X)[: , [2, 4]] # Select only  $x_1^2$ 
    ↪and  $x_2^2$ 

# Fit the logistic regression model to the square-only transformed data
model_circle_square = LogisticRegression()
model_circle_square.fit(X_circle_square, Y)

# Coefficients for the square-only features
```

```
print("Coefficients:", model_circle_square.coef_)
print("Intercept:", model_circle_square.intercept_)
```


Coefficients: [[4.91410958 4.97630742]]

Intercept: [-6.45841785]

f) Plot the decision boundary using the code below.

```
[7]: # create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
meshData = np.c_[xx.ravel(), yy.ravel()]
fig, ax = plt.subplots()
A = model.predict_proba(meshData)[:, 1].reshape(xx.shape)
Z = model.predict(meshData).reshape(xx.shape)
ax.contourf(xx, yy, A, cmap="RdBu", vmin=0, vmax=1)
ax.axis('off')

# Plot also the training points
ax.scatter(X[:, 0], X[:, 1], c=Y, s=50, alpha=0.9)
plt.show()
```




g) Plot the XOR data. In this 2D space, the data is not linearly separable, but by introducing a new feature

$$x_3 = x_1 * x_2$$

(called an interaction term) we should be able to find a hyperplane that separates the data in 3D. Plot this new dataset in 3D.

```
[8]: from mpl_toolkits.mplot3d import Axes3D

X, Y = generate_xor_data()
ax = plt.axes(projection='3d')
plt.figure(figsize=(8, 6))
ax.scatter3D(X[:, 0], X[:, 1], X[:, 0]*X[:, 1], c=Y)
plt.show()
```



<Figure size 800x600 with 0 Axes>

h) Apply a logistic regression model using the interaction term. Plot the decision boundary.

```
[9]: poly = PolynomialFeatures(interaction_only=True)
lr = LogisticRegression(verbose=0)
model = make_pipeline(poly, lr).fit(X, Y)

# create a mesh to plot in
h = .02 # step size in the mesh
```

```

x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))
meshData = np.c_[xx.ravel(), yy.ravel()]

fig, ax = plt.subplots()
A = model.predict_proba(meshData)[: , 1].reshape(xx.shape)
Z = model.predict(meshData).reshape(xx.shape)
ax.contourf(xx, yy, A, cmap="RdBu", vmin=0, vmax=1)
ax.axis('off')

# Plot also the training points
ax.scatter(X[:, 0], X[:, 1], color=Y, s=50, alpha=0.9)
plt.show()

```



```

[10]: '''
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LogisticRegression
%matplotlib widget

```



```

for i in range(20000):
    for solver in ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag',
        ↪ 'saga']:
        X_transform = PolynomialFeatures(interaction_only=True,
        ↪ include_bias=False).fit_transform(X)
        model = LogisticRegression(verbose=0, solver=solver, random_state=i,
        ↪ max_iter=10000)
        model.fit(X_transform, Y)
        print(model.score(X_transform, Y))
        if model.score(X_transform, Y) > .75:
            print("random state = ", i)
            print("solver = ", solver)
            break

print(model.coef_)
print(model.intercept_)

xx, yy = np.meshgrid([x / 10 for x in range(-1, 11)], [x / 10 for x in
    ↪ range(-1, 11)])
z = - model.intercept_ / model.coef_[0][2] - model.coef_[0][0] * xx / model.
    ↪ coef_[0][2] - model.coef_[0][1] * yy / model.coef_[0][2]

ax = plt.axes(projection='3d')
ax.scatter3D(X[:, 0], X[:, 1], X[:, 0] * X[:, 1], c=Y)
ax.plot_surface(xx, yy, z, alpha=0.5)
plt.show()
'''

```

```

[10]: '\nimport numpy as np\nimport matplotlib.pyplot as plt\nfrom
mpl_toolkits.mplot3d import Axes3D\nfrom sklearn.preprocessing import
PolynomialFeatures\nfrom sklearn.linear_model import
LogisticRegression\n%matplotlib widget\nfor i in range(20000):\n    for solver
in ['\lbfgs', '\liblinear', '\newton-cg', '\newton-cholesky', '\sag',
'\saga']:\n        X_transform = PolynomialFeatures(interaction_only=True,
include_bias=False).fit_transform(X)\n        model =
LogisticRegression(verbose=0, solver=solver, random_state=i, max_iter=10000)\n
model.fit(X_transform, Y)\n        print(model.score(X_transform, Y))\n
if model.score(X_transform, Y) > .75:\n            print("random state = ", i)\n
print("solver = ", solver)\n
break\n\nprint(model.coef_)\nprint(model.intercept_)\n\nxx, yy = np.meshgrid([x
/ 10 for x in range(-1, 11)], [x / 10 for x in range(-1, 11)])\nz = -
model.intercept_ / model.coef_[0][2] - model.coef_[0][0] * xx /
model.coef_[0][2] - model.coef_[0][1] * yy / model.coef_[0][2]\n\nax =
plt.axes(projection='\3d')\nax.scatter3D(X[:, 0], X[:, 1], X[:, 0] * X[:,
1], c=Y)\nax.plot_surface(xx, yy, z, alpha=0.5)\nplt.show()\n'

```

i) Using the code below that generates 3 concentric circles, fit a logistic regression model to it

and plot the decision boundary.

```
[11]: t, _ = datasets.make_blobs(n_samples=1500, centers=centers, cluster_std=2,
                                random_state=0)

# CIRCLES
def generate_circles_data(t):
    def label(x):
        if x[0]**2 + x[1]**2 >= 2 and x[0]**2 + x[1]**2 < 8:
            return 1
        if x[0]**2 + x[1]**2 >= 8:
            return 2
        return 0
    # create some space between the classes
    X = np.array(list(filter(lambda x : (x[0]**2 + x[1]**2 < 1.8 or x[0]**2 +
↪x[1]**2 > 2.2) and (x[0]**2 + x[1]**2 < 7.8 or x[0]**2 + x[1]**2 > 8.2), t)))
    Y = np.array([label(x) for x in X])
    return X, Y

X, Y = generate_circles_data(t)

poly = PolynomialFeatures(2)
lr = LogisticRegression(verbose=2)
model = make_pipeline(poly, lr).fit(X, Y)

#
h = 0.02
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                     np.arange(y_min, y_max, h))

#
Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

#
plt.contourf(xx, yy, Z, alpha=0.8, cmap='viridis')
plt.scatter(X[:, 0], X[:, 1], c=Y, edgecolors='k', s=20, cmap='viridis')
plt.title("Decision Boundary for 3 Concentric Circles")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 21 M = 10

At X0 0 variables are exactly at the bounds

At iterate	0	f=	1.09861D+00	proj g =	1.56782D+00
At iterate	1	f=	8.95383D-01	proj g =	5.58383D-01
At iterate	2	f=	8.27195D-01	proj g =	3.32189D-01
At iterate	3	f=	7.96120D-01	proj g =	6.06694D-01
At iterate	4	f=	7.66577D-01	proj g =	1.62504D-01
At iterate	5	f=	7.38227D-01	proj g =	1.74906D-01
At iterate	6	f=	6.22098D-01	proj g =	1.59122D-01
At iterate	7	f=	4.78311D-01	proj g =	1.20393D-01
At iterate	8	f=	2.88632D-01	proj g =	2.29261D-01
At iterate	9	f=	2.73371D-01	proj g =	3.32278D-01
At iterate	10	f=	1.61027D-01	proj g =	1.86164D-01
At iterate	11	f=	1.31057D-01	proj g =	8.20086D-02
At iterate	12	f=	1.09857D-01	proj g =	4.31233D-02
At iterate	13	f=	9.54732D-02	proj g =	3.14860D-02
At iterate	14	f=	8.82113D-02	proj g =	3.19970D-02
At iterate	15	f=	8.65265D-02	proj g =	2.58420D-02
At iterate	16	f=	8.51986D-02	proj g =	7.71859D-03
At iterate	17	f=	8.49323D-02	proj g =	4.32702D-03
At iterate	18	f=	8.47892D-02	proj g =	2.54731D-03
At iterate	19	f=	8.45880D-02	proj g =	4.72692D-03
At iterate	20	f=	8.43495D-02	proj g =	9.76915D-03

At iterate	21	f=	8.39228D-02	proj g =	1.41123D-02
At iterate	22	f=	8.30751D-02	proj g =	1.33183D-02
At iterate	23	f=	8.27346D-02	proj g =	1.95171D-02
At iterate	24	f=	8.16237D-02	proj g =	1.19741D-02
At iterate	25	f=	7.96364D-02	proj g =	7.19066D-03
At iterate	26	f=	7.76664D-02	proj g =	1.85242D-02
At iterate	27	f=	7.46684D-02	proj g =	3.03654D-02
At iterate	28	f=	7.05634D-02	proj g =	5.30341D-02
At iterate	29	f=	6.54279D-02	proj g =	3.59637D-02
At iterate	30	f=	6.18421D-02	proj g =	2.28137D-02
At iterate	31	f=	5.85717D-02	proj g =	6.21250D-03
At iterate	32	f=	5.74600D-02	proj g =	5.64093D-03
At iterate	33	f=	5.58955D-02	proj g =	1.31517D-02
At iterate	34	f=	5.38386D-02	proj g =	1.53870D-02
At iterate	35	f=	5.32192D-02	proj g =	1.92596D-02
At iterate	36	f=	5.19024D-02	proj g =	3.39938D-03
At iterate	37	f=	5.17519D-02	proj g =	1.98737D-03
At iterate	38	f=	5.16150D-02	proj g =	1.95118D-03
At iterate	39	f=	5.15147D-02	proj g =	2.96761D-03
At iterate	40	f=	5.14626D-02	proj g =	2.59971D-03
At iterate	41	f=	5.13984D-02	proj g =	1.62724D-03
At iterate	42	f=	5.13370D-02	proj g =	9.99570D-04
At iterate	43	f=	5.12664D-02	proj g =	1.37624D-03
At iterate	44	f=	5.11763D-02	proj g =	1.60635D-03

At iterate	45	f=	5.10465D-02	proj g =	1.34433D-03
At iterate	46	f=	5.10047D-02	proj g =	4.95510D-03
At iterate	47	f=	5.08214D-02	proj g =	4.46364D-03
At iterate	48	f=	5.07167D-02	proj g =	1.08309D-03
At iterate	49	f=	5.06903D-02	proj g =	8.37843D-04
At iterate	50	f=	5.06685D-02	proj g =	4.53040D-04
At iterate	51	f=	5.06555D-02	proj g =	1.17846D-03
At iterate	52	f=	5.06408D-02	proj g =	7.58418D-04
At iterate	53	f=	5.06250D-02	proj g =	9.22698D-04
At iterate	54	f=	5.05228D-02	proj g =	1.71929D-03
At iterate	55	f=	5.03782D-02	proj g =	1.97601D-03
At iterate	56	f=	5.02844D-02	proj g =	2.05844D-03
At iterate	57	f=	5.01694D-02	proj g =	9.05956D-04
At iterate	58	f=	5.01232D-02	proj g =	7.69174D-04
At iterate	59	f=	5.01149D-02	proj g =	8.32515D-04
At iterate	60	f=	5.01087D-02	proj g =	1.84661D-04
At iterate	61	f=	5.01070D-02	proj g =	1.57103D-04
At iterate	62	f=	5.01059D-02	proj g =	4.65643D-04
At iterate	63	f=	5.01039D-02	proj g =	2.51792D-04
At iterate	64	f=	5.01027D-02	proj g =	1.24677D-04
At iterate	65	f=	5.01016D-02	proj g =	1.21590D-04
At iterate	66	f=	5.01008D-02	proj g =	9.65235D-05

* * *

Tit = total number of iterations
Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value


* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
21	66	75	1	0	0	9.652D-05	5.010D-02

F = 5.0100827208558775E-002

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

This problem is unconstrained.



1.2 Gradient Descent

Recall in Linear Regression we are trying to find the line

$$y = X\beta$$

that minimizes the sum of square distances between the predicted y and the y we observed in our dataset:

$$\mathcal{L}(\beta) = \|\mathbf{y} - X\beta\|^2$$


We were able to find a global minimum to this loss function but we will try to apply gradient descent to find that same solution.

a) Implement the loss function to complete the code and plot the loss as a function of β .

```
[12]: %matplotlib widget
from mpl_toolkits import mplot3d
import numpy as np
import matplotlib.pyplot as plt

beta = np.array([ 1 , .5 ])
xlin = -10.0 + 20.0 * np.random.random(100)
X = np.column_stack([np.ones((len(xlin), 1)), xlin])
y = beta[0]+(beta[1]*xlin)+np.random.randn(100)

fig, ax = plt.subplots()
ax.plot(xlin, y, 'ro', markersize=4)
ax.set_xlim(-10, 10)
ax.set_ylim(min(y), max(y))
plt.show()
```



```
[13]: b0 = np.arange(-5, 4, 0.1)
      b1 = np.arange(-5, 4, 0.1)
      b0, b1 = np.meshgrid(b0, b1)

      def loss(X, y, beta):
          predictions = X.dot(beta)
          errors = y - predictions
          return np.sum(errors ** 2)

      def get_cost(B0, B1):
          res = []
          for b0, b1 in zip(B0, B1):
              line = []
              for i in range(len(b0)):
                  beta = np.array([b0[i], b1[i]])
                  line.append(loss(X, y, beta))
              res.append(line)
          return np.array(res)

      cost = get_cost(b0, b1)
```




```

# Creating figure
fig = plt.figure(figsize =(14, 9))
ax = plt.axes(projection ='3d')
ax.set_xlim(-6, 4)
ax.set_xlabel(r'\beta_0')
ax.set_ylabel(r'\beta_1')
ax.set_ylim(-5, 4)
ax.set_zlim(0, 30000)

# Creating plot
ax.plot_surface(b0, b1, cost, alpha=.7)

# show plot
plt.show()

```



Since the loss is

$$\mathcal{L}(\beta) = \|\mathbf{y} - \mathbf{X}\beta\|^2 = \beta^T \mathbf{X}^T \mathbf{X} \beta - 2 \mathbf{X}^T \mathbf{y} + \mathbf{y}^T \mathbf{y}$$

The gradient is

$$\nabla_{\beta} \mathcal{L}(\beta) = 2\mathbf{X}^T \mathbf{X} \beta - 2\mathbf{X}^T \mathbf{y}$$

b) Implement the gradient function below and complete the gradient descent algorithm

```

[14]: import numpy as np
from PIL import Image as im
import matplotlib.pyplot as plt

TEMPFILE = "temp.png"

def snap(betas, losses):
    # Creating figure
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(20, -20)
    ax.set_xlim(-5, 4)
    ax.set_xlabel(r'$\beta_0$')
    ax.set_ylabel(r'$\beta_1$')
    ax.set_ylim(-5, 4)
    ax.set_zlim(0, 30000)

    # Creating plot
    ax.plot_surface(b0, b1, cost, color='b', alpha=.7)
    ax.plot(np.array(betas)[: ,0], np.array(betas)[: ,1], losses, 'o-', c='r',
    ↪markersize=10, zorder=10)
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def gradient(X, y, beta):
    return 2 * X.T.dot(X.dot(beta) - y)

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap(betas, losses))
        beta_hat = beta_hat - learning_rate * gradient(X, y, beta_hat)

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

    return np.array(betas), np.array(losses)

beta_start = np.array([-5, -2])
learning_rate = 0.0002 # try .0005
images = []

```

```

betas, losses = gradient_descent(X, y, beta_start, learning_rate, 10, images)

images[0].save(
    'gd.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```

c) Use the code above to create an animation of the linear model learned at every epoch.

```

[15]: def snap_model(beta):
    xplot = np.linspace(-10,10,50)
    yestplot = beta[0] + beta[1] * xplot
    fig, ax = plt.subplots()
    ax.plot(xplot, yestplot, 'b-', lw=2)
    ax.plot(xlin, y, 'ro', markersize=4)
    ax.set_xlim(-10, 10)
    ax.set_ylim(min(y), max(y))
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap_model(beta_hat))
        beta_hat = beta_hat - learning_rate * gradient(X, y, beta_hat)

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

    return np.array(betas), np.array(losses)

images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, 100, images)

images[0].save(
    'model.gif',
    optimize=False,
    save_all=True,

```

```

    append_images=images[1:],
    loop=0,
    duration=200
)

```

In logistic regression, the `loss` is the negative log-likelihood

$$l(\beta) = -\frac{1}{N} \sum_{i=1}^N y_i \log(\sigma(x_i \beta)) + (1 - y_i) \log(1 - \sigma(x_i \beta))$$

the gradient of which is:

$$\nabla_{\beta} l(\beta) = \frac{1}{N} \sum_{i=1}^N x_i (y_i - \sigma(x_i \beta))$$

d) Plot the loss as a function of β .

```

[16]: %matplotlib widget
from mpl_toolkits import mplot3d
import numpy as np
import matplotlib.pyplot as plt
import sklearn.datasets as datasets


centers = [[0, 0]]
t, _ = datasets.make_blobs(n_samples=100, centers=centers, cluster_std=2,
    ↪random_state=0)

# LINE
def generate_line_data():
    # create some space between the classes
    X = t
    Y = np.array([1 if x[0] - x[1] >= 0 else 0 for x in X])
    return X, Y

X, y = generate_line_data()

cs = np.array([x for x in 'gb'])
fig, ax = plt.subplots()
ax.scatter(X[:, 0], X[:, 1], color=cs[y].tolist(), s=50, alpha=0.9)
plt.show()

```



```
[17]: b0 = np.arange(-20, 20, 0.1)
      b1 = np.arange(-20, 20, 0.1)
      b0, b1 = np.meshgrid(b0, b1)

      def sigmoid(x):
          e = np.exp(x)
          return e / (1 + e)

      def loss(X, y, beta):
          z = np.dot(X, beta)
          pred_prob = sigmoid(z)
          return -np.mean(y * np.log(pred_prob + 1e-5) + (1 - y) * np.log(1 -
          ↪pred_prob + 1e-5))

      def get_cost(B0, B1):
          res = []
          for b0, b1 in zip(B0, B1):
              line = []
```

```

    for i in range(len(b0)):
        beta = np.array([b0[i], b1[i]])
        line.append(loss(X, y, beta))
    res.append(line)
    return np.array(res)


cost = get_cost(b0, b1)

# Creating figure
fig = plt.figure(figsize=(14, 9))
ax = plt.axes(projection='3d')
ax.set_xlim(-20, 20)
ax.set_xlabel(r'$\beta_0$')
ax.set_ylabel(r'$\beta_1$')
ax.set_ylim(-20, 20)
ax.set_zlim(0, 10)

# Creating plot
ax.plot_surface(b0, b1, cost, alpha=.7)

# show plot
plt.show()

```



e) Plot the loss at each iteration of the gradient descent algorithm.

```

[18]: import numpy as np
from PIL import Image as im
import matplotlib.pyplot as plt

TEMPFILE = "temp.png"

def snap(betas, losses):
    # Creating figure
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(10, 10)
    ax.set_xlabel(r'$\beta_0$')
    ax.set_ylabel(r'$\beta_1$')
    ax.set_ylim(-20, 20)
    ax.set_zlim(0, 10)

    # Creating plot
    ax.plot_surface(b0, b1, cost, color='b', alpha=.7)
    ax.plot(np.array(betas)[: ,0], np.array(betas)[: ,1], losses, 'o-', c='r',
    ↪markersize=10, zorder=10)
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def gradient(X, y, beta):
    z = X.dot(beta)
    y_pred = sigmoid(z)
    gradient = X.T.dot(y_pred - y)
    return gradient

def gradient_descent(X, y, beta_hat, learning_rate, epochs, images):
    losses = [loss(X, y, beta_hat)]
    betas = [beta_hat]

    for _ in range(epochs):
        images.append(snap(betas, losses))
        beta_hat = beta_hat - learning_rate * gradient(X, y, beta_hat)

        losses.append(loss(X, y, beta_hat))
        betas.append(beta_hat)

```

```

    return np.array(betas), np.array(losses)

beta_start = np.array([-5, -2])
learning_rate = 0.1
images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, 10, images)

images[0].save(
    'gd_logit.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```

f) Create an animation of the logistic regression fit at every epoch.

```

[19]: beta_start = np.array([-5, -2])
learning_rate = 0.1
images = []
betas, losses = gradient_descent(X, y, beta_start, learning_rate, 10, images)

images[0].save(
    'gd_logit.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```

g) Modify the above code to evaluate the gradient on a random batch of the data. Overlay the true loss curve and the approximation of the loss in your animation.

```

[20]: import numpy as np
from sklearn.utils import shuffle

def mini_batch_gradient_descent(X, y, beta_hat, learning_rate, epochs,
    batch_size, images):
    full_losses = [loss(X, y, beta_hat)]
    batch_losses = []
    betas = [beta_hat.copy()]

    for epoch in range(epochs):
        X, y = shuffle(X, y) # Shuffle the dataset

```



```

    for i in range(0, X.shape[0], batch_size):
        X_batch = X[i:i + batch_size]
        y_batch = y[i:i + batch_size]

        # Calculate gradient using the batch
        grad = gradient(X_batch, y_batch, beta_hat)
        beta_hat -= learning_rate * grad
        betas.append(beta_hat.copy())

        # Calculate and store batch loss for visualization
        batch_current_loss = loss(X_batch, y_batch, beta_hat)
        batch_losses.append(batch_current_loss)

        # Calculate full dataset loss after the update for true trajectory
        ↪ visualization
        full_current_loss = loss(X, y, beta_hat)
        full_losses.append(full_current_loss)

        # Snapshot current state for animation
        images.append(snap(betas, full_losses, batch_losses))

    return np.array(betas), np.array(full_losses), np.array(batch_losses)

def snap(betas, full_losses, batch_losses):
    min_length = min(len(betas), len(full_losses), len(batch_losses))
    betas = np.array(betas)[:min_length]
    full_losses = np.array(full_losses)[:min_length]
    batch_losses = np.array(batch_losses)[:min_length]

    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(10, 10)
    ax.set_xlabel(r'$\beta_0$')
    ax.set_ylabel(r'$\beta_1$')
    ax.set_ylim(-20, 20)
    ax.set_zlim(0, max(max(full_losses), max(batch_losses), 1))

    ax.plot_surface(b0, b1, cost, color='blue', alpha=.7)
    ax.plot(betas[:, 0], betas[:, 1], full_losses, 'o-', color='red', ↪
    ↪ markersize=10, zorder=10)
    ax.plot(betas[:, 0], betas[:, 1], batch_losses, 'o-', color='green', ↪
    ↪ markersize=5, zorder=5)

    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

```

```

beta_start = np.array([-5, -2], dtype=np.float64)
learning_rate = 0.1
images = []

betas, full_losses, batch_losses = mini_batch_gradient_descent(X, y,
    ↪beta_start, learning_rate, 10, 10, images)

images[0].save(
    'mini_batch_gd_logit.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```

h) Below is a sandox where you can get intuition about how to tune gradient descent parameters:

```

[21]: import numpy as np
from PIL import Image as im
import matplotlib.pyplot as plt

TEMPFILE = "temp.png"

def snap(x, y, pts, losses, grad):
    fig = plt.figure(figsize=(14, 9))
    ax = plt.axes(projection='3d')
    ax.view_init(20, -20)
    ax.plot_surface(x, y, loss(np.array([x, y])), color='r', alpha=.4)
    ax.plot(np.array(pts)[: ,0], np.array(pts)[: ,1], losses, 'o-', c='b',
    ↪markersize=10, zorder=10)
    ax.plot(np.array(pts)[-1,0], np.array(pts)[-1,1], -1, 'o-', c='b', alpha=.
    ↪5, markersize=7, zorder=10)

    # Plot Gradient Vector
    X, Y, Z = [pts[-1][0]], [pts[-1][1]], [-1]
    U, V, W = [-grad[0]], [-grad[1]], [0]
    ax.quiver(X, Y, Z, U, V, W, color='g')
    fig.savefig(TEMPFILE)
    plt.close()
    return im.fromarray(np.asarray(im.open(TEMPFILE)))

def loss(x):
    return np.sin(sum(x**2)) # change this

```

```

def gradient(x):
    return 2 * x * np.cos(sum(x**2)) # change this

def gradient_descent(x, y, init, learning_rate, epochs):
    images, losses, pts = [], [loss(init)], [init]
    for _ in range(epochs):
        grad = gradient(init)
        images.append(snap(x, y, pts, losses, grad))
        init = init - learning_rate * grad
        losses.append(loss(init))
        pts.append(init)
    return images

init = np.array([-0.5, -0.5]) # change this
learning_rate = 1.394 # change this
x, y = np.meshgrid(np.arange(-2, 2, 0.1), np.arange(-2, 2, 0.1)) # change this
images = gradient_descent(x, y, init, learning_rate, 12)

images[0].save(
    'gradient_descent.gif',
    optimize=False,
    save_all=True,
    append_images=images[1:],
    loop=0,
    duration=500
)

```